

Exercise

AR: Auto-Regressive, **MA:** Moving Average

```
In [30]: # Imports
import warnings
from statsmodels.tools.sm_exceptions import ConvergenceWarning
warnings.simplefilter('ignore', ConvergenceWarning)
import requests
import numpy as np
from io import BytesIO
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
from statsmodels.tsa.stattools import ccf, adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.arima_process import arma_generate_sample
from statsmodels.tsa.holtwinters import ExponentialSmoothing
np.random.seed(0) # For reproducibility
```

```
In [9]: def ts_plot(x,y,title,xlabel="Time",ylabel="Value"):
    plt.figure(figsize=(10,5))
    plt.plot(x, y)
    plt.title(title)
    plt.xlabel(xlabel)
    plt.ylabel(ylabel)
    plt.grid(alpha=0.3);
```

```
In [10]: def acf_pacf_plot(ts,lags=30):
    fig, ax = plt.subplots(1,2, figsize=(15, 5))
    plot_acf(ts, lags=lags, ax=ax[0])
    plot_pacf(ts, lags=lags, ax=ax[1])
    plt.show()
```

```
In [11]: def mse(observations, estimates):
    difference = observations - estimates
    sq_diff = difference ** 2
    mse = np.mean(sq_diff)
    return mse
```

Exercise

- Load the two time series `arma_tsi` by running the code below.

```
In [12]: # Load the first time series
response = requests.get("https://zenodo.org/records/10951538/files/arma_ts3.npz?download=1")
response.raise_for_status()
ts = np.load(BytesIO(response.content))['signal']
print(len(ts))
```

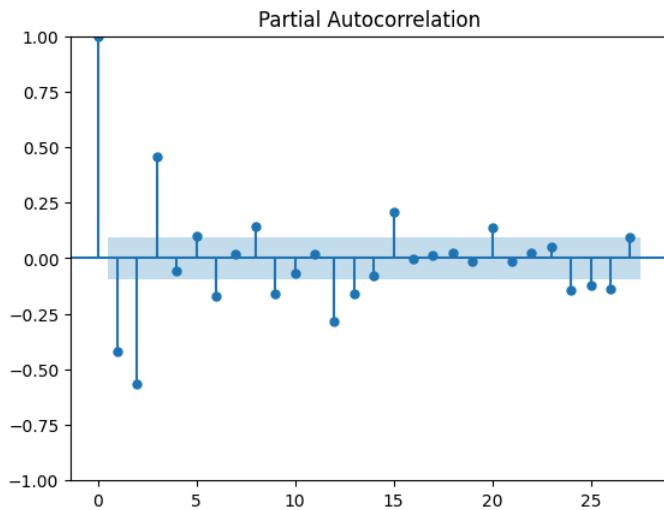
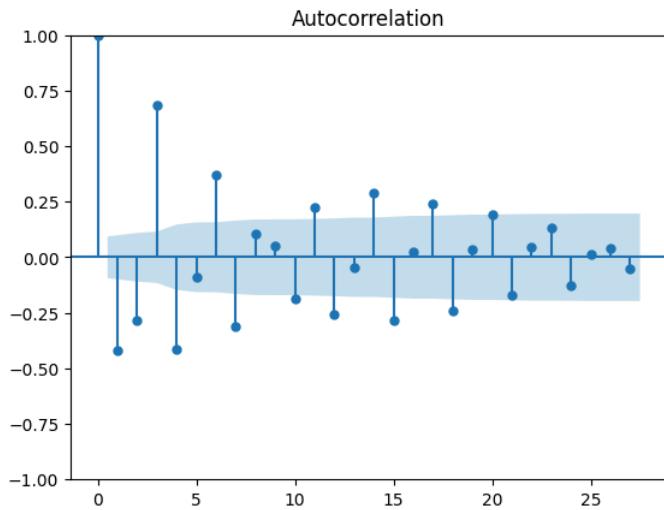
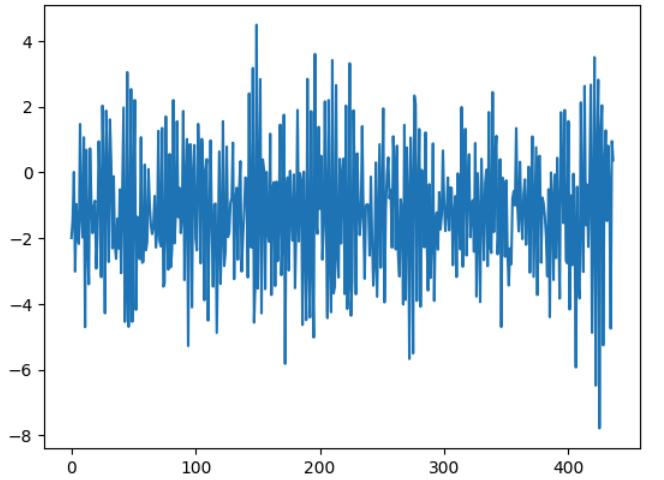
479

For each time series

- Split the time series in train and test.
 - Use the last 30 values as test for the first time series.
 - Use the last 100 as test for the second time series.
- Make the time series stationary.
- Determine the order p of an AR model.
- Compute the prediction of the test data with the AR(p) model.
- Determine the order q of a MA model.
- Compute the prediction of the test data with the MA(q) model.

```
In [13]: time = range(len(ts))
```

```
In [14]: train = ts[:-30]
test = ts[-30:]
s = 12
ts_diff = np.diff(ts)
ts_sdiff = ts[s:] - ts[:-s]
train_diff = ts_diff[:-30] #np.diff(train)
train_sdiff = ts_sdiff[:-30] #train[s:] - train[:-s]
test_diff = ts_diff[-30:]
test_sdiff = ts_sdiff[-30:]
# import pandas as pd
# train_sdiff = pd.DataFrame(train).diff(12).dropna()
# fig,ax=plt.subplots((1,3))
plt.plot(train_sdiff);
plot_acf(train_sdiff);
plot_pacf(train_sdiff);
```



Implementing ADF test

In [15]: `adfuller(train_sdiff)`

Out[15]: (-6.63876415493461,
 $5.477142531888239 \times 10^{-9}$,
14,
422,
{'1%': -3.44594128742536,
'5%': -2.868413360220551,
'10%': -2.570431271085555},
1343.7496673604292)

we see that `sdiff` is stationary

In [16]: `model_ARMA = ARIMA(train_sdiff, order=(3,0,0))
model_ARMA_fit = model_ARMA.fit() # Fit the model`

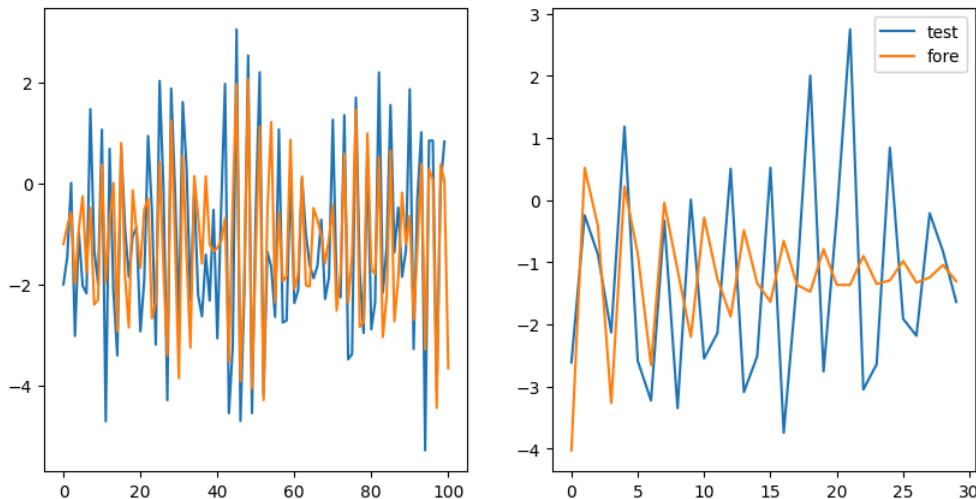
```
fig, axs = plt.subplots(1,2,figsize=(10,5))
axs[0].plot(train_sdiff[:100],label='train')
axs[0].plot(model_ARMA_fit.predict(0,100),label='pred')
axs[1].plot(test_sdiff,label='test')
```

```

    axs[1].plot(model_ARMA_fit.forecast(len(test_sdiff)),label='fore')
    plt.legend()

```

Out[16]: <matplotlib.legend.Legend at 0x1bfdd19ce90>



Auto-ARIMA

In [17]: `from pmдарима import auto_arima`

In [18]: `auto_arima(train_sdiff).summary(#, seasonal=True, m=7)`

Out[18]: SARIMAX Results

Dep. Variable:	y	No. Observations:	437			
Model:	SARIMAX(3, 0, 5)	Log Likelihood	-724.221			
Date:	Wed, 17 Dec 2025	AIC	1468.442			
Time:	13:24:06	BIC	1509.241			
Sample:	0	HQIC	1484.542			
	- 437					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
intercept	-0.4493	0.141	-3.193	0.001	-0.725	-0.174
ar.L1	-0.3248	0.051	-6.317	0.000	-0.426	-0.224
ar.L2	0.2419	0.054	4.454	0.000	0.135	0.348
ar.L3	0.7090	0.046	15.276	0.000	0.618	0.800
ma.L1	-0.0166	0.062	-0.267	0.790	-0.138	0.105
ma.L2	-0.5543	0.061	-9.054	0.000	-0.674	-0.434
ma.L3	-0.0984	0.067	-1.474	0.141	-0.229	0.032
ma.L4	0.1608	0.052	3.093	0.002	0.059	0.263
ma.L5	-0.4500	0.048	-9.302	0.000	-0.545	-0.355
sigma2	1.6336	0.125	13.058	0.000	1.388	1.879
Ljung-Box (L1) (Q):	0.34	Jarque-Bera (JB):	0.85			
Prob(Q):	0.56	Prob(JB):	0.65			
Heteroskedasticity (H):	1.23	Skew:	-0.06			
Prob(H) (two-sided):	0.21	Kurtosis:	2.83			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

Note that auto-arima considers in-sample performance

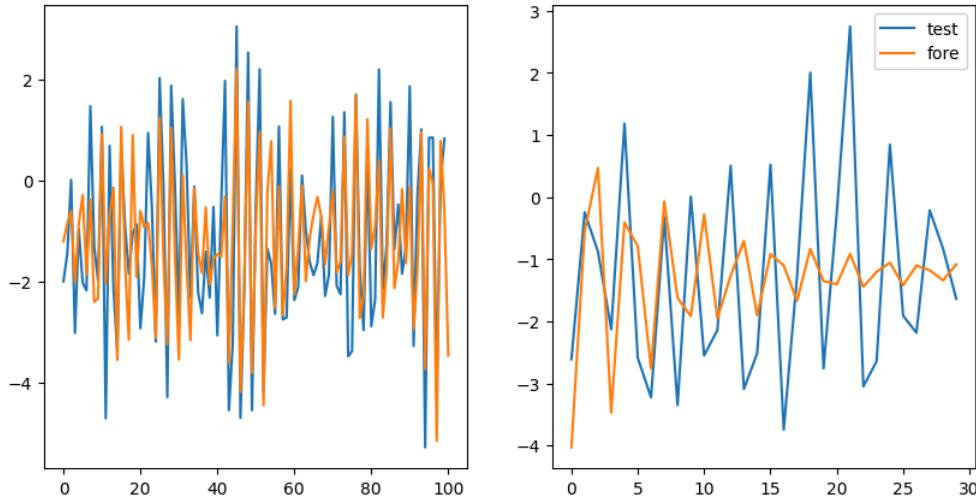
```

In [45]: model_ARMA_auto = ARIMA(train_sdiff, order=(3,0,5))
model_ARMA_auto_fit = model_ARMA_auto.fit() # Fit the model

fig, axs = plt.subplots(1,2,figsize=(10,5))
axs[0].plot(train_sdiff[:100],label='train')
axs[0].plot(model_ARMA_auto_fit.predict(0,100),label='pred')
axs[1].plot(test_sdiff,label='test')
axs[1].plot(model_ARMA_auto_fit.forecast(len(test_sdiff)),label='fore')
plt.legend()

```

Out[45]: <matplotlib.legend.Legend at 0x1bfdcb0bef0>



Grid Search

```
In [20]: %%time
best_RMSE = np.inf
for p in range(1,4):
    for d in [0]:
        for q in range(1,4):
            try:
                model = ARIMA(train_sdiff, order=(p,d,q)) # ARIMA with d=0 is equivalent to ARMA
                model_ARMA_grid_fit = model.fit()
                forecast_sdiff = model_ARMA_grid_fit.forecast(len(test_sdiff))
            except:
                pass
            else:
                MAE = abs((test_sdiff-forecast_sdiff)).mean() # not suitable for data containing 0s
                RMSE = mse(test_sdiff, forecast_sdiff)**0.5
                AIC = model_ARMA_grid_fit.aic
                BIC = model_ARMA_grid_fit.bic
                if RMSE < best_RMSE:
                    best_model = (p,d,q)
                    best_RMSE = RMSE
            print('ARIMA(%d,%d,%d) -> MAE:%.3f, RMSE:%.3f, AIC:%.1f, BIC:%.1f' % (p,d,q, MAE, RMSE, AIC, BIC))
print('Best model by RMSE (Forecast):',best_model)

ARIMA(1,0,1) -> MAE:1.412, RMSE:1.655, AIC:1719.6, BIC:1735.9
ARIMA(1,0,2) -> MAE:1.414, RMSE:1.660, AIC:1656.3, BIC:1676.7
C:\Users\serhan.aydin\Documents\my_docs\venv\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. Using zeros as starting parameters.
  warn('Non-invertible starting MA parameters found.')
ARIMA(1,0,3) -> MAE:1.400, RMSE:1.651, AIC:1574.0, BIC:1598.5
ARIMA(2,0,1) -> MAE:1.338, RMSE:1.601, AIC:1511.4, BIC:1531.8
ARIMA(2,0,2) -> MAE:1.347, RMSE:1.612, AIC:1505.4, BIC:1529.9
ARIMA(2,0,3) -> MAE:1.495, RMSE:1.750, AIC:1492.0, BIC:1520.6
ARIMA(3,0,1) -> MAE:1.437, RMSE:1.719, AIC:1499.8, BIC:1524.3
ARIMA(3,0,2) -> MAE:1.408, RMSE:1.677, AIC:1497.8, BIC:1526.4
ARIMA(3,0,3) -> MAE:1.484, RMSE:1.743, AIC:1493.1, BIC:1525.7
Best model by RMSE (Forecast): (2, 0, 1)
CPU times: total: 1.17 s
Wall time: 1.25 s

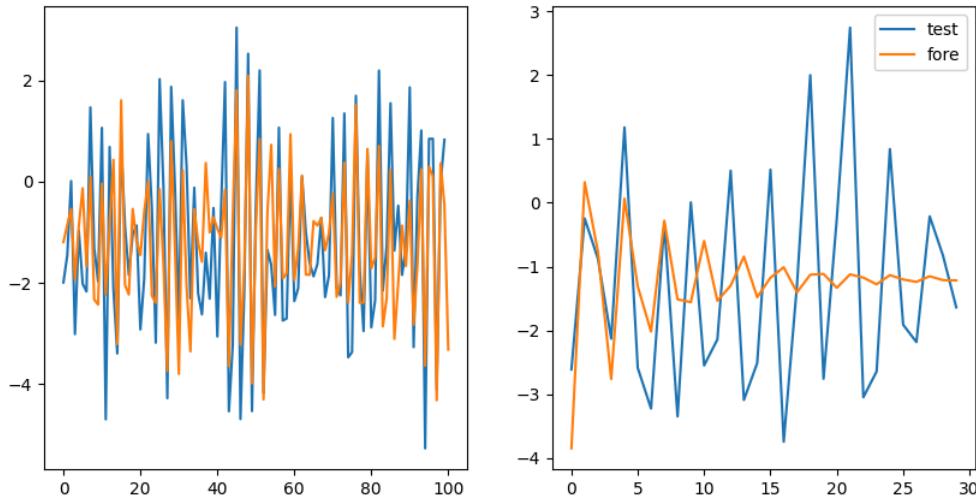
Do not use MAPE on:
```

- Differenced data
- Seasonally differenced data
- Log-returns
- Residual-like series

```
In [21]: model_ARMA_grid = ARIMA(train_sdiff, order=best_model)
model_ARMA_grid_fit = model_ARMA_grid.fit() # Fit the model

fig, axs = plt.subplots(1,2, figsize=(10,5))
axs[0].plot(train_sdiff[:100],label='train')
axs[0].plot(model_ARMA_grid_fit.predict(0,100),label='pred')
axs[1].plot(test_sdiff,label='test')
axs[1].plot(model_ARMA_grid_fit.forecast(len(test_sdiff)),label='fore')
plt.legend()
```

```
Out[21]: <matplotlib.legend.Legend at 0x1bfe047e210>
```



```
In [46]: from statsmodels.tools.eval_measures import mse
print('MSE train (model_ARMA_fit) %s:' % str(model_ARMA.order), mse(model_ARMA_fit.predict(),train_sdiff))
print('MSE train (model_ARMA_auto_fit) %s:' % str(model_ARMA_auto.order), mse(model_ARMA_auto_fit.predict(),train_sdiff))
print('MSE train (model_ARMA_grid_fit) %s:' % str(model_ARMA_grid.order), mse(model_ARMA_grid_fit.predict(),train_sdiff))
print('\n')
print('MSE test (model_ARMA_fit) %s:' % str(model_ARMA.order), mse(model_ARMA_fit.forecast(len(test_sdiff)),test_sdiff))
print('MSE test (model_ARMA_auto_fit) %s:' % str(model_ARMA_auto.order), mse(model_ARMA_auto_fit.forecast(len(test_sdiff)),test_sdiff))
print('MSE test (model_ARMA_grid_fit) %s:' % str(model_ARMA_grid.order), mse(model_ARMA_grid_fit.forecast(len(test_sdiff)),test_sdiff))

MSE train (model_ARMA_fit) (3, 0, 0): 1.7685451058780566
MSE train (model_ARMA_auto_fit) (3, 0, 5): 1.594820466280704
MSE train (model_ARMA_grid_fit) (2, 0, 1): 1.8128786923346902

MSE test (model_ARMA_fit) (3, 0, 0): 3.3326137214316542
MSE test (model_ARMA_auto_fit) (3, 0, 5): 2.5988643812573473
MSE test (model_ARMA_grid_fit) (2, 0, 1): 2.563046523485501
```

Discussion on the results:

- Auto-arima is based on in-sample
- This can explain why it is best in-sample but 2nd best in out-of-sample
- Visual check resulted in a relatively poorer performance

Reintegrate original series

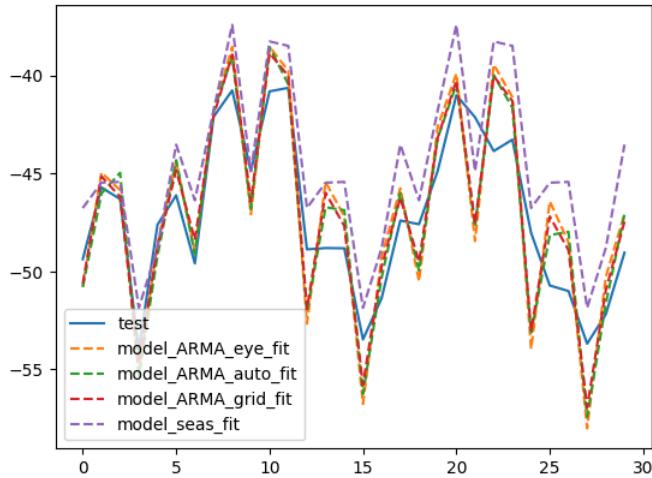
```
In [50]: # Reversing the seasonal differencing
test_fore_ARMA = np.empty(len(test))
test_fore_ARMA_auto = np.empty(len(test))
test_fore_ARMA_grid = np.empty(len(test))
test_fore_seas = np.empty(len(test))

test_fore_ARMA[:12] = train[-12:] + model_ARMA_fit.forecast(len(test_sdiff))[:12]
test_fore_ARMA_auto[:12] = train[-12:] + model_ARMA_auto_fit.forecast(len(test_sdiff))[:12]
test_fore_ARMA_grid[:12] = train[-12:] + model_ARMA_grid_fit.forecast(len(test_sdiff))[:12]
test_fore_seas[:12] = train[-12:]

for i in range(12, len(test)):
    test_fore_seas[i] = test_fore_seas[i-12]
    test_fore_ARMA[i] = test_fore_ARMA[i-12] + model_ARMA_fit.forecast(len(test_sdiff))[i]
    test_fore_ARMA_auto[i] = test_fore_ARMA_auto[i-12] + model_ARMA_auto_fit.forecast(len(test_sdiff))[i]
    test_fore_ARMA_grid[i] = test_fore_ARMA_grid[i-12] + model_ARMA_grid_fit.forecast(len(test_sdiff))[i]

plt.plot(test, label='test')
plt.plot(test_fore_ARMA,'--', label='model_ARMA_eye_fit')
plt.plot(test_fore_ARMA_auto,'--', label='model_ARMA_auto_fit')
plt.plot(test_fore_ARMA_grid,'--', label='model_ARMA_grid_fit')
plt.plot(test_fore_seas,'--', label='model_seas_fit')
plt.legend()
```

```
Out[50]: <matplotlib.legend.Legend at 0x1bfe64e67e0>
```



```
In [25]: from statsmodels.tools.eval_measures import mse
print('MSE test (model_ARMA_fit) %s:' % str(model_ARMA.order), mse(test_fore_ARMA,test))
print('MSE test (model_ARMA_auto_fit) %s:' % str(model_ARMA_auto.order), mse(test_fore_ARMA_auto,test))
print('MSE test (model_ARMA_grid_fit) %s:' % str(model_ARMA_grid.order), mse(test_fore_ARMA_grid,test))

MSE test (model_ARMA_fit) (3, 0, 0): 7.853957943219815
MSE test (model_ARMA_autoarima_fit) (3, 0, 5): 5.672869974333564
MSE test (model_ARMA_grid_fit) (2, 0, 1): 5.211031373991924
```

Alternatively we can use SARIMA fit/forecasting. But we have to tell the model how we achieved stationarity.

```
In [52]: model_SARIMA_grid_fit = SARIMAX(train, order=(2,0,1), seasonal_order=(0,1,0,12)).fit()
model_SARIMA_auto_fit = SARIMAX(train, order=(5,0,2), seasonal_order=(0,1,0,12)).fit()
model_SARIMA_eye_fit = SARIMAX(train, order=(3,0,0), seasonal_order=(0,1,0,12)).fit()

plt.plot(test, label='test')
plt.plot(model_SARIMA_grid_fit.forecast(len(test)), '--', label='model_SARIMA_grid_fore')
plt.plot(model_SARIMA_auto_fit.forecast(len(test)), '--', label='model_SARIMA_auto_fore')
plt.plot(model_SARIMA_eye_fit.forecast(len(test)), '--', label='model_SARIMA_eye_fore')
# plt.plot(test_fore_seas, '--', label='model_seas_fit')
plt.legend()
```

Out[52]: <matplotlib.legend.Legend at 0x1bfe65b6cc0>

